

Predicting user behaviour based on e-commerce data

Author: Tino Merl

Table of Contents

- [Introduction and planned action](#)
 - [CRISP-DM](#)
- [Business Understanding](#)
- [Data Understanding](#)
 - [Describing the files](#)
 - [Describing the columns](#)
 - [Focus dataset](#)
- [Data Preparation](#)
 - [Plotting the data by groups](#)
- [Modeling](#)
 - [Selection](#)
 - [Model Preparations](#)
 - [Fitting](#)
- [Evaluation](#)
- [Conclusion](#)

Introduction and planned action

This is an assignment for the module applied programming in the summerterm of 2020 at the *FOM Hochschule für Oekonomie & Management* at the study center in cologne. In this work a dataset which contains user behaviour data from an e-commerce system will be analysed. The dataset can be found on kaggle named *Retailrocket recommender system dataset*.^[1] It contains four individual files. The main focus of this notebook indeed lies on one dataset as stated in [Focus dataset](#). The goal is to predict user behaviour. This can be done in two ways. The data can be engineered and the users can be grouped in a way to predict whether a user belongs to the buying users or not. It can also be done by using markov chains to calculate the probability of a user buying or not. Since the latter is way more complex and is beyond the scope of this analysis therefore it will be tried to engineer the data and predict buyers by new features. This whole analysis will be done via the CRISP-DM Process.

CRISP-DM

CRISP-DM stands for **C**ross **I**ndustry **S**tandard **P**rocess for **D**ata **M**ining. It is a framework which describes the most common steps a machine learning analysis and model building should undertake. The steps are the following six.

1. Business Understanding
2. Data Understanding
3. Data Preparation
4. Modeling
5. Evaluation
6. Deployment

They may be listed in a sequential manner, but there is a lot of back and forth between the steps. Especially between business understanding and data understanding, business understanding and evaluation such as data preparation and modeling. Figure 1 illustrates the circular nature of the process.[2]



Figure 1: CRISP-DM diagram by statistik-dresden.de[3]

Since this is an assignment and not a real business case the last step, the deployment, will be left out. The analysis will therefore end after step number five: evaluation.

Business Understanding

This focus of the first step is to lay out and understand the objectives and requirements from a business perspective. The concrete tasks will also be defined here as it can be defined as a planning phase.[2] When looking at the data the main question about it is what kind of insights can be generated from it? Another one which can be described by the following: Is it possible to find out whether a user will buy an item or not? This will be the main question of the oncoming analysis. But before commencing with any analysis the tasks will be defined and explained them in a detailed way.

Tasks

1. *Explorative analysis of the data in data understanding*
2. *Cluster unusual user behaviour and filter them out*
3. *How can user behaviour be predicted?*

1. Explorative analysis of the data in data understanding

This is a really huge dataset. It contains more than 2 million rows.[1] To even grasp what information is in these files and what hidden gems may be in them the data exploration must be thoroughly. It is arguably that this is clearly part of the data understanding chapter, but for this assignment it was made more explicit.

2. Cluster unusual user behaviour and filter them out

Filtering out the unusual user behaviour can go two ways. In some analysis this unusual behaviour could be the main focus as to find out why there are outliers and how to get these back on track. To catch these outliers and put them "back on the right track" may be the goal of some customer analysis. But for this task the goal is to try filtering them out and find out if this has any impact on the modeling stage. There are many users in the dataset who didn't buy anything or only had a very limited numbers of transactions. Therefore is a lot of noise in the dataset. This noise could be prone to overfitting.[1]

3. How can user behaviour be predicted?

The most important question of this analysis for user behaviour is whether or not the behaviour can be predicted based on historical data. This is the whole goal of this assignment. To find out whether or not the behaviour, especially whether or not someone will buy something, can be predicted. For this case different models will be tested and a look at their respective prediction accuracy will be took.

Data Understanding

The goal of this chapter is trying to get familiar with the data and to identify problems or first insights from it.[2] The first step to do so is by an explorative analysis of the dataset, which also creates the foundation for the following chapter: the data preparation. Since this dataset has a usability score of 8.8 kaggle and has also a lot of context describing the dataset it's kaggle page will mostly be cited.[1]

Describing the files

As already mentioned the dataset consists of four files. One of them is splitted into two parts since their respective filesize goes above the kaggle size limit. Therefore two of the files can be counted as one.[1] They will also be described as one file.

events.csv

A file containing the user behaviour. E.g. if a user viewed an item or added it to their cart with the corresponding item and visitorid and a respective timestamp.[1]

item_properties_part1.csv and item_properties_part2.csv

Two files containing the item properties. This refers to the property and the value of the item as their respective itemid and a timestamp. These files contain weekly snapshots of the item properties. If the property or the value of an item differs between these weekly snapshots a new row for this item will be added to the files. The tricky part of this dataset is how the columns `property` and `value` refer to each other. Every value in these two columns was hashed. Except for the entries `available` and `categoryid`. When the column `property` has the entry `available` the column `value` refers to the availability of the item. Containing 1 for available and 0 for not. However if `property` contains the entry `categoryid` `value` refers to the item's category referencing the `category_tree.csv` file. All other values were stemmed and hashed. However the numeric values are marked with a n.[1] The description of the dataset is cited as followed:

All words in text values were normalized (stemming procedure: <https://en.wikipedia.org/wiki/Stemming> (<https://en.wikipedia.org/wiki/Stemming>)) and hashed, numbers were processed as above, e.g. text "Hello world 2017!" will become "24214 44214 n2017.000"[1]

category_tree.csv

This file specifies a categoryid and their respective parentid.[\[1\]](#)

Describing the columns

The following descriptions are also citations from the kaggle dataset.

events.csv

- timestamp {int} -- the time, when an event occurred, in milliseconds since 01-01-1970.
- visitorid {int} -- unique identifier of the visitor.
- event {string} -- type of the event {"view", "addtocart", "transaction"}.
- itemid {int} -- unique identifier of the item.
- transactionid {int} -- unique identifier of the transaction (non empty only for transaction event type).[\[1\]](#)

item_properties_part1.csv & item_properties_part2.csv

- timestamp {int} -- snapshot creation time (unix timestamp in milliseconds).
- itemid {int} -- unique id of the item.
- property {str} -- property of the item. All of them had been hashed excluding "categoryid" and "available".
- value {str} -- property value of the item.[\[1\]](#)

category_tree.csv

- categoryid {int} -- unique identifier of the category.
- parentid {int} -- identifier of the parent category. It's empty, if parent doesn't exist.[\[1\]](#)

Focus dataset

After looking at the descriptions of the files and their respective columns it seems, that not all files may be relevant to the question if a user buys an item or not. Only one file, to be precise the **events.csv**, seems to hold the relevant information. The reasoning for this assumption is the following. Starting with the item properties files. They are currently hashed and their column `value` contains a mixture of different values. They range from hashed text, hashed numbers, a combination of them and normal text or numbers. A change in value or pricing may lead to a user buying but since the values are all hashed it's very difficult to make out a difference in prices in this file. Also to get a grasp about what is inside this dataset and whether or not it could be applied to the main question of this assignment is out of scope. Since the item properties were already excluded to cancel any analysis on **category_tree.csv** is only logical. It doesn't hold any direct relation to the **events.csv** and can only be used by utilizing the item properties. Therefore the now following quantitative and qualitative analysis will focus solely on the **events.csv** file.

In [1]:



```
# Loading of the needed packages
import datetime
import warnings

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.preprocessing import LabelEncoder

# Surpress any warnings
warnings.filterwarnings("ignore")

# Setting seed
np.random.seed(28)
```

The first step of the quantative analysis is to look at the dataframe with descriptive and explorative methods. The findings and their meanings will be described and why they may be added to the list of the data preparations. It will also be tried to go into more detail so there is less to write in the part of data preparation.

In [2]:



```
# Loading dataset
eventsDf = pd.read_csv("./data/events.csv")
eventsDf.sample(10)
```

Out[2]:

	timestamp	visitorid	event	itemid	transactionid
629976	1438573436544	895853	view	432161	NaN
348678	1434570301279	229157	view	422669	NaN
937150	1439964270842	217570	view	368947	NaN
2316508	1436751804165	1015325	view	333515	NaN
2123125	1436004758561	770911	view	321219	NaN
2218502	1436402747169	689710	view	204494	NaN
1097762	1440778593222	749691	view	298883	NaN
2275467	1436523377814	965534	view	227279	NaN
1835554	1432131850263	625059	view	343947	NaN
172391	1433905443612	672582	view	53707	NaN

By looking at our sample there are two initial findings. The first one was clear from the description of the files and their columns. `timestamp` is a unix timestamp and should be casted to a datetime object for better interpretability. Finding number two is that `transactionid` seems to be casted as a float object. This is a normal behaviour since the majority of the column looks to be empty and whenever a column contains NaN values and numerics it will be casted as float. Just to make sure this assumption is right a look at the datatypes of `eventsDf` will be taken.

In [3]:



```
eventsDf.dtypes
```

Out[3]:

```
timestamp      int64
visitorid      int64
event          object
itemid         int64
transactionid  float64
dtype: object
```

This reassures both assumptions made above. Now built-in describe method will be used to analyze the quantitative metrics for this dataset.

In [4]:



```
eventsDf.describe(include="all")
```

Out[4]:

	timestamp	visitorid	event	itemid	transactionid
count	2.756101e+06	2.756101e+06	2756101	2.756101e+06	22457.000000
unique	NaN	NaN	3	NaN	NaN
top	NaN	NaN	view	NaN	NaN
freq	NaN	NaN	2664312	NaN	NaN
mean	1.436424e+12	7.019229e+05	NaN	2.349225e+05	8826.497796
std	3.366312e+09	4.056875e+05	NaN	1.341954e+05	5098.996290
min	1.430622e+12	0.000000e+00	NaN	3.000000e+00	0.000000
25%	1.433478e+12	3.505660e+05	NaN	1.181200e+05	4411.000000
50%	1.436453e+12	7.020600e+05	NaN	2.360670e+05	8813.000000
75%	1.439225e+12	1.053437e+06	NaN	3.507150e+05	13224.000000
max	1.442545e+12	1.407579e+06	NaN	4.668670e+05	17671.000000

Except for the already mentioned empty columns in `transactionid` there seems to be nothing out of ordinary. None of the other columns has any indications that any cell may be empty since their counts are all on the same level. Since the `timestamp` is still in unix format and `visitorid` as `itemid` are id columns it is hard to make any assumptions based on descriptive statistics. Nevertheless it is noteworthy that the most common entry in the column `event` is "view" with 2664312 occurrences. Since there are only two other unique values that means there are only 91789 left for the other two levels of `event`. This should be mentioned because the dataset therefore has a massive overrepresentation of the top value while the other two values "addtocart" and "transaction", which are more or less the ones that are interesting, are less likely to be found.

As of next the frequencies in the data will be plotted. It was already mentioned that `visitorid` and `itemid` are id columns. It is therefore hard to make any assumptions for this columns. Nevertheless they will be included in the plotting to find out if there may be a pattern.

In [5]:



```
fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20, 15))

ax[0, 0].hist(eventsDf["event"])
ax[0, 0].set_title("events")
ax[0, 1].hist(eventsDf["transactionid"].fillna(0))
ax[0, 1].set_title("transactionids")
ax[0, 2].hist(eventsDf["itemid"].dropna())
ax[0, 2].set_title("itemids")
ax[1, 0].hist(eventsDf["timestamp"].dropna())
ax[1, 0].set_title("timestamps")
ax[1, 1].hist(eventsDf["visitorid"].dropna())
ax[1, 1].set_title("visitorids")

plt.draw()
```

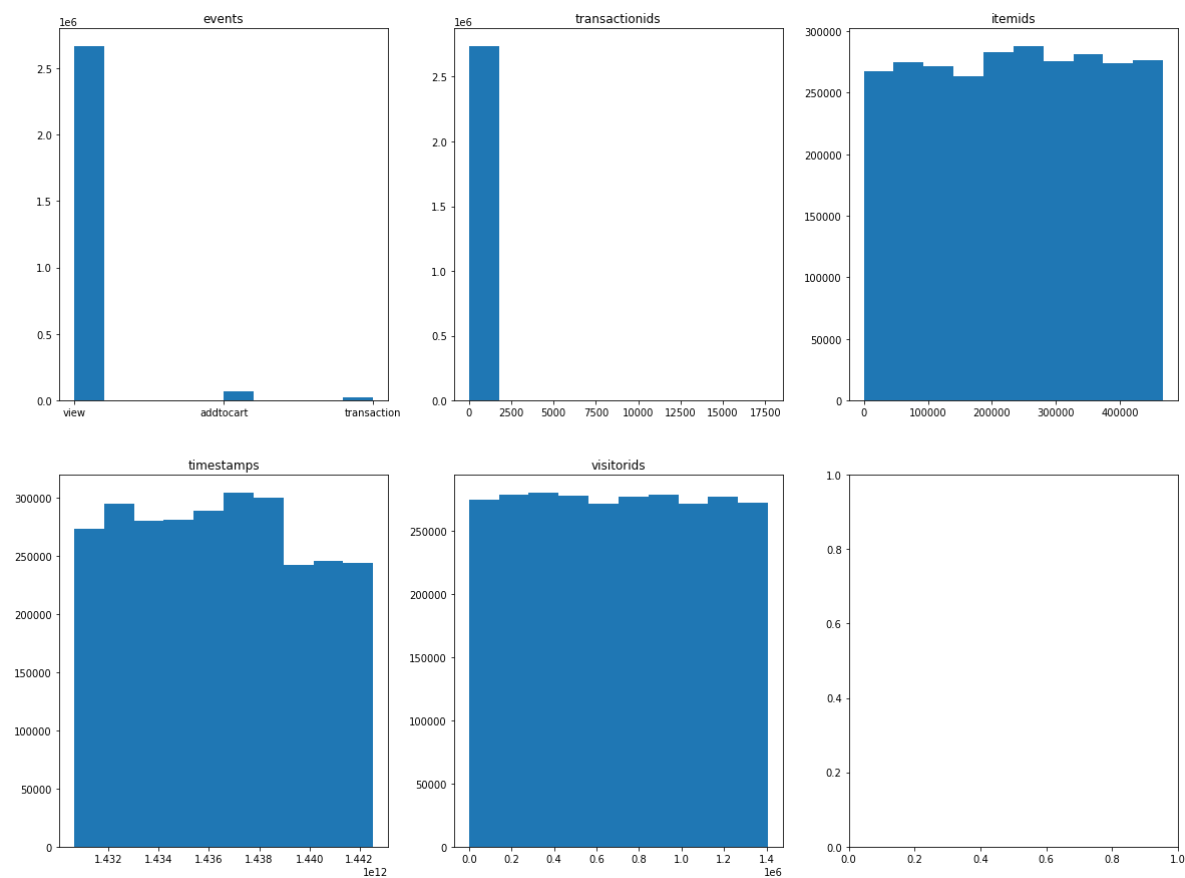


Figure 2: Occurrences of values for every column.

This further concludes the assumption that most of the users don't have any transactions. This can be clearly seen in the top left histogram named event. Also the most occurrences of transactionids seem to be at zero. visitorid and itemid don't seem to have any particular pattern. Only timestamp has more occurrences in the lower timestamps. This could be a sudden drop in users.

Since the buying of an item is underrepresented the next question will be how many users actually bought something. These users will be saved in a new variable called buyingUsers. The percentage of buying users will then be calculated. Therefore the userIds which made a transaction will first be filtered out.

In [6]:



```
# Filter out unique visitorids that made a transaction
buyingUsers = pd.DataFrame(
    eventsDf[eventsDf.event == "transaction"]["visitorid"]
    .copy()
    .drop_duplicates(keep="first")
    .reset_index(drop=True)
)

# Calculate percentage of users buying
percBuying = buyingUsers.shape[0] / eventsDf["visitorid"].unique().shape[0]

print(f"Unique number of users who made a transaction: {buyingUsers.shape[0]}")
print(
    f"Percentage of unique users who made a transaction: {round(percBuying * 100, 2)}%"
)
```

Unique number of users who made a transaction: 11719

Percentage of unique users who made a transaction: 0.83%

The number of visitors who made a transaction and bought something is 11719. This number alone doesn't say a lot. But compared to the number of all visitors it will give a better outlook of how many users made a transaction. The percentage of buying users is less than 1% it is 0.83%. That is a very small fraction the focus should be laid upon to identify potential buyers. So the further analysis must be conducted carefully as to not overfit the model. The last thing to look at now is whether or not any insights can be generated from the timestamp. A useful idea would be to calculate sessions per user. If a session timer is set to a certain threshold, e.g. 30 minutes between two actions, the actions a user took can be grouped by sessions. To find out if the column is of any use for this question a user who bought an item is randomly selected and a look at their customer journey is taken.

In [7]:



```

sampleUser = (
    eventsDf[eventsDf["event"] == "transaction"]["visitorid"].sample(1).tolist()[0]
)
sampleDf = (
    eventsDf[eventsDf["visitorid"] == sampleUser]
    .sort_values("timestamp")
    .reset_index(drop=True)
)

sampleDf["timestamp"] = pd.to_datetime(sampleDf["timestamp"], unit="ms")
sampleDf

```

Out[7]:

	timestamp	visitorid	event	itemid	transactionid
0	2015-05-05 15:10:44.676	227091	view	168491	NaN
1	2015-05-05 15:11:55.893	227091	view	168491	NaN
2	2015-05-05 15:12:28.817	227091	view	461710	NaN
3	2015-05-05 15:12:52.830	227091	view	200507	NaN
4	2015-05-05 15:13:57.618	227091	view	168491	NaN
...
687	2015-05-28 18:32:10.454	227091	addtocart	150173	NaN
688	2015-05-28 18:33:23.880	227091	transaction	150173	15947.0
689	2015-05-28 18:42:34.655	227091	view	277082	NaN
690	2015-05-28 18:53:13.440	227091	view	440560	NaN
691	2015-05-28 18:58:36.187	227091	view	68607	NaN

692 rows × 5 columns

The timestamps look promising and there is already a case visible where the user bought an item. In row 688 the user made a transaction. A closer look at the last 10 actions that led to this transaction will now be taken.

In [8]:



sampleDf[678:689]

Out[8]:

	timestamp	visitorid	event	itemid	transactionid
678	2015-05-28 18:10:20.416	227091	view	82278	NaN
679	2015-05-28 18:20:13.101	227091	view	306045	NaN
680	2015-05-28 18:24:44.999	227091	view	157065	NaN
681	2015-05-28 18:24:53.411	227091	view	157065	NaN
682	2015-05-28 18:25:01.848	227091	view	130231	NaN
683	2015-05-28 18:25:12.267	227091	view	184284	NaN
684	2015-05-28 18:30:33.089	227091	view	150173	NaN
685	2015-05-28 18:31:04.687	227091	view	310720	NaN
686	2015-05-28 18:31:52.884	227091	view	60759	NaN
687	2015-05-28 18:32:10.454	227091	addtocart	150173	NaN
688	2015-05-28 18:33:23.880	227091	transaction	150173	15947.0

One can clearly see, that the timestamp is in a usable format for this case. However the data preparations will be a bit more difficult based on the sheer number of entries.

Data Preparation

The focus of this chapter will be to construct the final dataset. Here the feature engineering and cleaning of the data will take place.[\[2\]](#) This work is defined by the following tasks.

- *Convert timestamp to datetime.*
- *Calculate total session time per user.*
- *Calculate average session time per user.*
- *Calculate number of sessions a user had*
- *Filter out unusual user behaviour.*

The Features to engineer are the following:

- Number of items a user viewed.
- Numer of unique items a user viewed.
- Numer of times a user added an item to a cart.
- Number of times a user bought something.
- Number of sessions a user had.
- Average session time per user.
- Total Session time per user.

To calculate whether an action belongs to a session or not a certain threshold will need to be set. For this example it will be at 30 minutes. So to speak whenever two actions of a user are more than 30 minutes apart a new session "will be started". This will need to be calculated for every user. But the first step is to cast the

timestamp column to datetime. The dataframe will also be sorted first by `visitorid` and then by `timestamp` in an ascending order. The sorting is needed so it'll be easier to identify if an action by a certain user belongs to the same session or a new one.

In [9]:



```
# Cast timestamp to datetime and sort by visitorid and timestamp
eventsDf["timestamp"] = pd.to_datetime(eventsDf["timestamp"], unit="ms")
eventsDf = eventsDf.sort_values(["visitorid", "timestamp"]).reset_index(drop=True)
eventsDf.head(10)
```

Out[9]:

	timestamp	visitorid	event	itemid	transactionid
0	2015-09-11 20:49:49.439	0	view	285930	NaN
1	2015-09-11 20:52:39.591	0	view	357564	NaN
2	2015-09-11 20:55:17.175	0	view	67045	NaN
3	2015-08-13 17:46:06.444	1	view	72028	NaN
4	2015-08-07 17:51:44.567	2	view	325215	NaN
5	2015-08-07 17:53:33.790	2	view	325215	NaN
6	2015-08-07 17:56:52.664	2	view	259884	NaN
7	2015-08-07 18:01:08.920	2	view	216305	NaN
8	2015-08-07 18:08:25.669	2	view	342816	NaN
9	2015-08-07 18:17:24.375	2	view	342816	NaN

After doing so it is now possible to calculate the time difference between each action. What is done afterwards is taking the current timestamp of every row and subtract the timestamp from the row before to get the duration. It is no problem that the time difference between different users is also calculated. Reason why is because a new column named `sameUser` will also be added. This one compares the `visitorid` of the current row with the one above. If it is the same it returns `True` otherwise `False`. The last step is then to add a row named `sessionlimit`. This one is `True` when the actions are within the threshold and `False` if not.

In [10]:



```
# Calculate time differences
eventsDf["duration"] = eventsDf["timestamp"] - eventsDf["timestamp"].shift()

# Look up if the visitorid is the same as in the row above
eventsDf["sameUser"] = eventsDf["visitorid"].eq(eventsDf["visitorid"].shift())

# Calculate if a duration between two actions is within our set treshold of 30
# minutes
limit = datetime.timedelta(minutes=30)
eventsDf["sessionlimit"] = eventsDf["duration"].apply(
    lambda x: True if x <= limit else False
)

eventsDf.head(10)
```

Out[10]:

	timestamp	visitorid	event	itemid	transactionid	duration	sameUser	sessionlimit
0	2015-09-11 20:49:49.439	0	view	285930	NaN	NaT	False	False
1	2015-09-11 20:52:39.591	0	view	357564	NaN	0 days 00:02:50.152000	True	True
2	2015-09-11 20:55:17.175	0	view	67045	NaN	0 days 00:02:37.584000	True	True
3	2015-08-13 17:46:06.444	1	view	72028	NaN	-30 days +20:50:49.269000	False	True
4	2015-08-07 17:51:44.567	2	view	325215	NaN	-6 days +00:05:38.123000	False	True
5	2015-08-07 17:53:33.790	2	view	325215	NaN	0 days 00:01:49.223000	True	True
6	2015-08-07 17:56:52.664	2	view	259884	NaN	0 days 00:03:18.874000	True	True
7	2015-08-07 18:01:08.920	2	view	216305	NaN	0 days 00:04:16.256000	True	True
8	2015-08-07 18:08:25.669	2	view	342816	NaN	0 days 00:07:16.749000	True	True
9	2015-08-07 18:17:24.375	2	view	342816	NaN	0 days 00:08:58.706000	True	True

It is now finished to look up whether actions belong to the same session or a new one. A column named `session` can now be added. It will be a counter. It starts at 1 for every user. If the column `sameUser` is `True` and `sessionlimit` is `False` 1 will be added to `session`. Afterwards a new column `sessionid` which is a merge from `visitorid` and `session` will be created. This is used to create a unique column per session per user.



In [11]:

```

sessionList = []
session = 1

# Iterate over every row from the columns sameUser and sessionLimit
for r in zip(eventsDf["sameUser"], eventsDf["sessionLimit"]):

    # If a user is not the same session counter will be reset
    if r[0] == False:
        session = 1

    # If a user is the same but actions are over the timelimit a new
    # session will begin
    elif r[0] == True and r[1] == False:
        session += 1
    sessionList.append(session)

# Add everything to the dataframe and create a new column
eventsDf["session"] = sessionList
eventsDf["sessionid"] = (
    eventsDf["visitorid"].astype(str) + "_" + eventsDf["session"].astype(str)
)

eventsDf.head(10)

```

Out[11]:

	timestamp	visitorid	event	itemid	transactionid	duration	sameUser	sessionlimit
0	2015-09-11 20:49:49.439	0	view	285930	NaN	NaT	False	False
1	2015-09-11 20:52:39.591	0	view	357564	NaN	0 days 00:02:50.152000	True	True
2	2015-09-11 20:55:17.175	0	view	67045	NaN	0 days 00:02:37.584000	True	True
3	2015-08-13 17:46:06.444	1	view	72028	NaN	-30 days +20:50:49.269000	False	True
4	2015-08-07 17:51:44.567	2	view	325215	NaN	-6 days +00:05:38.123000	False	True
5	2015-08-07 17:53:33.790	2	view	325215	NaN	0 days 00:01:49.223000	True	True
6	2015-08-07 17:56:52.664	2	view	259884	NaN	0 days 00:03:18.874000	True	True
7	2015-08-07 18:01:08.920	2	view	216305	NaN	0 days 00:04:16.256000	True	True
8	2015-08-07 18:08:25.669	2	view	342816	NaN	0 days 00:07:16.749000	True	True
9	2015-08-07 18:17:24.375	2	view	342816	NaN	0 days 00:08:58.706000	True	True

Since the dataframe now also contains the information about sessions a session duration can now be calculated. To do so the timestamp of the first action of a session will be subtracted from the last action's timestamp. The parameters `subset` and `keep` of pandas built in method `drop_duplicates` will be utilized to first get the first action of a session and afterwards the last one. They will then be merged in a new dataframe

and the timestamps will be subtracted. The new column `total_sessionsduration` will then be a `timedelta` object. Since transformations with such a `timedelta` object are computational heavy it will be casted into seconds as an integer. This makes it faster to group the columns by `visitorid` to get the sum of all sessions a user had.

In [12]:

```
# Getting the starting timestamp of every session
sessionsStart = eventsDf.drop_duplicates(
    subset=["visitorid", "sessionid"], keep="first"
)[["timestamp", "visitorid", "sessionid"]]

# Getting the ending timestamp of every session
sessionsEnd = eventsDf.drop_duplicates(subset=["visitorid", "sessionid"], keep="last")[
    ["timestamp", "visitorid", "sessionid"]
]

# Merging them to get a new dataframe
sessionsDuration = sessionsStart.merge(
    sessionsEnd, on=["visitorid", "sessionid"], how="left"
)

# Calculating the duration between start and end and also casting it to seconds
sessionsDuration["total_sessionsduration"] = (
    sessionsDuration["timestamp_y"] - sessionsDuration["timestamp_x"]
).dt.seconds

# Group by visitorid to get the total duration of all sessions
sessionsDuration = (
    sessionsDuration.drop(columns=["timestamp_x", "timestamp_y", "sessionid"])
    .groupby(["visitorid"], as_index=False)
    .sum()
)
sessionsDuration.head(10)
```

Out[12]:

	visitorid	total_sessionsduration
0	0	327
1	1	0
2	2	1753
3	3	0
4	4	0
5	5	0
6	6	1015
7	7	186
8	8	0
9	9	0

The next step is to get the total number of sessions a user had. This is a fairly simple operation. The dataframe just needs to be reduced to the columns `visitorid` and `session` and afterwards it can be grouped by `visitorid` calculating the max of `sessions` per user.

In [13]:



```
# Reducing to the columns visitorid and session and getting the max value for session
sessionsNums = (
    eventsDf[["visitorid", "session"]]
    .groupby(["visitorid"], as_index=False)
    .max()
    .rename(columns={"session": "num_sessions"})
)
sessionsNums.head(10)
```

Out[13]:

	visitorid	num_sessions
0	0	1
1	1	1
2	2	1
3	3	1
4	4	1
5	5	1
6	6	3
7	7	2
8	8	1
9	9	1

The preparations in the area of user sessions are now finished. Therefore it is possible to move on and calculate the number of unique item views. To do so firstly a subset of the dataframe is created to have only the event view remaining, drop the duplicates with the subset of of visitorid and itemid keeping the first occurrences. It will then be grouped by visitorid counting the number of entries.

In [14]:



```
# Getting unique views on items
uniqueViews = (
    eventsDf[eventsDf["event"] == "view"]
    .drop_duplicates(subset=["visitorid", "itemid"], keep="first")
    ["visitorid", "itemid"]
)
.sort_values(["visitorid"])

# Grouping unique item Views
uniqueViews = (
    uniqueViews.groupby(["visitorid"], as_index=False)
    .count()
    .rename(columns={"itemid": "unique_item_views"})
)

uniqueViews.head(10)
```

Out[14]:

	visitorid	unique_item_views
0	0	3
1	1	1
2	2	4
3	3	1
4	4	1
5	5	1
6	6	2
7	7	3
8	8	1
9	9	1

The point of merging everything together is almost reached. But before doing so it is needed to add a new column named `bought` to the dataframe `buyingUsers`. By merging it with `eventsDf` a new dataframe called `procEventsDf` is created. Shorthanded for *processed events dataframe*. The reason behind doing so is not needing to reread the whole dataset when a false data preparation step should be undone. Next is to fill the empty values in `bought` with 0 in the new dataframe as to identify users who didn't buy anything. After that three new columns are added which are one hot encoded from the column `events`. This will help calculating the total number of actions for every type of action per user. To get these sums group by is used again and the sums of the new made columns are added.



In [15]:

```
# Adding a new column to the dataframe of buying users
buyingUsers["bought"] = 1

# Merging the two dataframes
procEventsDf = pd.merge(eventsDf, buyingUsers, on="visitorid", how="left")

# Replacing NaNs
procEventsDf["bought"].fillna(0, inplace=True)

# Get number of actions for every possible event
procEventsDf = pd.get_dummies(procEventsDf, columns=["event"], prefix="items")
procEventsDf = (
    procEventsDf[
        ["visitorid", "bought", "items_view", "items_addtocart", "items_transaction"]
    ]
    .groupby(["visitorid", "bought"], as_index=False)
    .sum()
)

# Cast everything to integer
procEventsDf = procEventsDf.astype("int64")
procEventsDf.sort_values(["visitorid"]).head(10)
```

Out[15]:

	visitorid	bought	items_view	items_addtocart	items_transaction
0	0	0	3	0	0
1	1	0	1	0	0
2	2	0	8	0	0
3	3	0	1	0	0
4	4	0	1	0	0
5	5	0	1	0	0
6	6	0	5	1	0
7	7	0	3	0	0
8	8	0	1	0	0
9	9	0	1	0	0

The final step is to merge all the feature dataframes to `procEventsDf` and calculate the mean duration of sessions per user. After this the feature engineering is finished and some manual outlier detection will take place.



In [16]:

```
# Merge with unique views
procEventsDf = procEventsDf.merge(uniqueViews, on="visitorid", how="left")

# Merge with session duration
procEventsDf = procEventsDf.merge(sessionsDuration, on="visitorid", how="left")

# Merge with number of sessions
procEventsDf = procEventsDf.merge(sessionsNums, on="visitorid", how="left")

# Calculate average session time
procEventsDf["mean_sessionsduration"] = (
    procEventsDf["total_sessionsduration"] / procEventsDf["num_sessions"]
)
procEventsDf.head(10)
```

Out[16]:

	visitorid	bought	items_view	items_addtocart	items_transaction	unique_item_views	total_s
0	0	0	3	0	0	3.0	
1	1	0	1	0	0	1.0	
2	2	0	8	0	0	4.0	
3	3	0	1	0	0	1.0	
4	4	0	1	0	0	1.0	
5	5	0	1	0	0	1.0	
6	6	0	5	1	0	2.0	
7	7	0	3	0	0	3.0	
8	8	0	1	0	0	1.0	
9	9	0	1	0	0	1.0	

What can be seen now is a dataframe with every visitorid, whether they bought or not, their number of items view, their number of items added to cart, their number of transactions, their number of unique item views, their total duration of all their sessions, their number of sessions and their mean session duration. With these values it should be possible to find out about the unusual user behaviour. But what first strikes as odd is that the column `unique_item_views` now looks like it was casted as a float instead of integer. As mentioned in the beginning this is usually a sign that there are empty rows present. Which means there are users who have never viewed any item. Before going further ahead this should be checked.

In [17]:



```
procEventsDf[procEventsDf["unique_item_views"].isna()]
```

Out[17]:

	visitorid	bought	items_view	items_addtocart	items_transaction	unique_item_views
1159	1159	0	0	1	0	NaN
1399	1399	0	0	2	0	NaN
1480	1480	0	0	1	0	NaN
1742	1742	0	0	1	0	NaN
2277	2277	0	0	1	0	NaN
...
1406820	1406820	0	0	1	0	NaN
1407059	1407059	0	0	1	0	NaN
1407319	1407319	0	0	1	0	NaN
1407398	1407398	1	0	1	1	NaN
1407515	1407515	0	0	1	0	NaN

3401 rows × 9 columns



There are really users who don't have any entry in the column `unique_item_views`. And they also have a zero in `items_view`. But at least one of them has bought an item. Since this could still be an error in the data wrangling this should be checked in the raw data. Therefore some of these ids are used as filters in the original dataframe.

In [18]:



```
eventsDf[eventsDf["visitorid"].isin([1159, 1399, 1406820, 1407398])]
```

Out[18]:

	timestamp	visitorid	event	itemid	transactionid	duration	sameUser
2112	2015-07-27 22:59:58.750	1159	addtocart	337777	NaN	9 days 18:07:55.810000	False
2539	2015-05-10 22:07:31.393	1399	addtocart	136523	NaN	-65 days +06:33:28.489000	False
2540	2015-05-10 22:07:31.708	1399	addtocart	136523	NaN	0 days 00:00:00.315000	True
2754703	2015-07-09 21:38:04.767	1406820	addtocart	377888	NaN	20 days 04:13:26.148000	False
2755789	2015-07-07 03:34:38.082	1407398	addtocart	218917	NaN	53 days 05:14:07.181000	False
2755790	2015-07-07 03:47:15.969	1407398	transaction	218917	10009.0	0 days 00:12:37.887000	True



The assumptions made above were right. There are users inside this dataset who didn't ever see any item which is really odd. This is the first unusual behaviour that should be treated carefully. Shouldn't users see at least one item before proceeding to buy? But to make sure this doesn't make any problems the NaNs will be replaced with zero and the column is casted to int.

In [19]:



```
procEventsDf["unique_item_views"].fillna(0, inplace=True)
procEventsDf = procEventsDf.astype({"unique_item_views": "int64"})
procEventsDf[procEventsDf["unique_item_views"].isna()]
```

Out[19]:

visitorid	bought	items_view	items_addtocart	items_transaction	unique_item_views	total_se
-----------	--------	------------	-----------------	-------------------	-------------------	----------

To find out more unusual behaviour a closer look at the columns `items_view` and `unique_item_views` will be taken and they will be plotted them as separate groups.

Plotting the data by groups

For this the data will be parted into two groups. Users who bought an item and those who don't. The first pairs of plots will be about unusual behaviour. The second pair is used to find out if there's any linearity or clear distinctions in the newly engineered dataframe.

In [20]:



```
fig, ax = plt.subplots(ncols=2, nrows=2, figsize=(20, 15))

ax[0, 0].hist(procEventsDf[procEventsDf["bought"] == 0]["items_view"])
ax[0, 0].set_title("Non Buyers total number of item views")
ax[0, 1].hist(procEventsDf[procEventsDf["bought"] == 1]["items_view"])
ax[0, 1].set_title("Buyers total number of item views")
ax[1, 0].hist(procEventsDf[procEventsDf["bought"] == 0]["unique_item_views"])
ax[1, 0].set_title("Non Buyers unique number of item views")
ax[1, 1].hist(procEventsDf[procEventsDf["bought"] == 1]["unique_item_views"])
ax[1, 1].set_title("Buyers unique number of item views")

plt.draw()
```




Figure 3.1: Plotting buying and non buying users by their number of total item views and unique item views

Most Users seem to have a number of below 100 views. It also doesn't matter whether the user is buying or not and also the unique number of views. The graphs will now be recreated with limited ranges. The focus will therefore be on the total views and with different thresholds.



In [21]:

```
fig, ax = plt.subplots(ncols=2, nrows=4, figsize=(20, 20))

ax[0, 0].hist(
    procEventsDf[(procEventsDf["bought"] == 0) & (procEventsDf["items_view"] >= 3)][
        "items_view"
    ]
)
ax[0, 0].set_title("Non Buyers total number of item views if >= 3")

ax[0, 1].hist(
    procEventsDf[(procEventsDf["bought"] == 1) & (procEventsDf["items_view"] >= 3)][
        "items_view"
    ]
)
ax[0, 1].set_title("Buyers total number of item views if >= 3")

ax[1, 0].hist(
    procEventsDf[(procEventsDf["bought"] == 0) & (procEventsDf["items_view"] >= 3)][
        "unique_item_views"
    ]
)
ax[1, 0].set_title("Non Buyers unique number of item views if >= 3")

ax[1, 1].hist(
    procEventsDf[(procEventsDf["bought"] == 1) & (procEventsDf["items_view"] >= 3)][
        "unique_item_views"
    ]
)
ax[1, 1].set_title("Buyers unique number of item views if >= 3")

ax[2, 0].hist(
    procEventsDf[(procEventsDf["bought"] == 0) & (procEventsDf["items_view"] <= 100)][
        "items_view"
    ]
)
ax[2, 0].set_title("Non Buyers total number of item views if <= 100")

ax[2, 1].hist(
    procEventsDf[(procEventsDf["bought"] == 1) & (procEventsDf["items_view"] <= 100)][
        "items_view"
    ]
)
ax[2, 1].set_title("Buyers total number of item views if <= 100")

ax[3, 0].hist(
    procEventsDf[(procEventsDf["bought"] == 0) & (procEventsDf["items_view"] <= 100)][
        "unique_item_views"
    ]
)
ax[3, 0].set_title("Non Buyers unique number of item views if <= 100")

ax[3, 1].hist(
    procEventsDf[(procEventsDf["bought"] == 1) & (procEventsDf["items_view"] <= 100)][
        "unique_item_views"
    ]
)
ax[3, 1].set_title("Buyers unique number of item views if <= 100")

plt.draw()
```




Figure 3.2: Plotting buying and non buying users by their number of total item views and unique item views with different ranges

These graphs indicate, that the most outliers seem to be the users which have more than 100 views. The first distinction between the buying users and those who don't can also be seen. While users that don't buy any items have the most number of views in the range from 0 to 10 views and a small fraction in the area from 10 to 20 the buying users tend to be a little bit more spread in the higher numbers. On the graphs on the right side it can be seen, that users who made a transaction also have a strong base in the area of 0 to 10 views but they also are more spreading towards the higher numbers. As a last graphical analysis the two levels above will now be combined.



In [22]:

```
fig, ax = plt.subplots(ncols=2, nrows=2, figsize=(20, 15))

ax[0, 0].hist(
    procEventsDf[
        (procEventsDf["bought"] == 0)
        & (procEventsDf["items_view"] >= 3)
        & (procEventsDf["items_view"] <= 100)
    ]["items_view"]
)
ax[0, 0].set_title("Non Buyers total number of item views if >= 3 & <= 100")
ax[0, 1].hist(
    procEventsDf[
        (procEventsDf["bought"] == 1)
        & (procEventsDf["items_view"] >= 3)
        & (procEventsDf["items_view"] <= 100)
    ]["items_view"]
)
ax[0, 1].set_title("Buyers total number of item views if => 3 & <= 100")
ax[1, 0].hist(
    procEventsDf[
        (procEventsDf["bought"] == 0)
        & (procEventsDf["items_view"] >= 3)
        & (procEventsDf["items_view"] <= 100)
    ]["unique_item_views"]
)
ax[1, 0].set_title("Non Buyers unique number of item views if => 3 & <= 100")
ax[1, 1].hist(
    procEventsDf[
        (procEventsDf["bought"] == 1)
        & (procEventsDf["items_view"] >= 3)
        & (procEventsDf["items_view"] <= 100)
    ]["unique_item_views"]
)
ax[1, 1].set_title("Buyers unique number of item views if => 3 & <= 100")

plt.draw()
```






Figure 3.3: Plotting buying and non buying users by their number of total item views and unique item views with combined ranges.

This hardens the assumptions made above. The most density of users is in the range of 3 to 100 views. This filtering was made completely manual. While it's safe to say this is a way to do it, the question arises whether or not these outliers could be detected automatically. The reasoning behind this is that in the future user behaviour might change. If certain thresholds are set manually, they could be outdated in the near future. In reality, it would be tried to filter out these outliers by anomaly detection methods like a classifier. For the sake of this assignment, and since the data won't change, this will be restricted to manually filtering out the data.

As it was already found out in the graphs above, there are definitely outliers which may make the dataset prone to overfitting. The question to answer now is if there's any linearity in the dataset and if there are any clear distinctions between the groups. To do so, the pairplot method from seaborn is used, and the dataset is reduced to the range found out above.

In [23]:



```
sns.pairplot(  
    procEventsDf[  
        (procEventsDf["items_view"] >= 3) & (procEventsDf["items_view"] <= 100)  
    ],  
    x_vars=[  
        "items_view",  
        "unique_item_views",  
        "items_transaction",  
        "total_sessionsduration",  
        "mean_sessionsduration",  
        "num_sessions",  
    ],  
    y_vars=[  
        "items_view",  
        "unique_item_views",  
        "items_transaction",  
        "total_sessionsduration",  
        "mean_sessionsduration",  
        "num_sessions",  
    ],  
    hue="bought",  
)
```

Out[23]:

<seaborn.axisgrid.PairGrid at 0x7fd3cf5dac50>





Figure 4: Pairplot for the engineered features to guess the linearity.

In the scatterplots one is able to make out some linearity. The most common and obvious linearity is between `unique_item_views` and `item_views`. This should be no surprise. However what surprises is that there seems to be no real difference between the two groups. But one thing different is, that the group of non buying users tend spread well throughout the lower ranks of `item_views`. A not too new finding considering the manual outlier detection. The only clear distinction can be made in the combination of `total_sessionsduration` and `num_sessions`. The non buying users tend to have more unique sessions than buying users but tend to have a lower `total_sessionsduration`. In every other combination the groups are overlapping.

Modeling

Selection

As for the model selection the range is limited to classifiers. Some of the more common classifiers are

- Logistic regression
- Gaussian naive bayes
- Random forest classifier
- Support vector classification
- Decision tree classifier

- Nearest Neighbors[4]

Since the features are not really independent and there are high correlations between the features a logistic regression can not be used as a classifier.[5] This is supported by the correlation matrix for our features.

In [24]:

```
corr = procEventsDf.drop(columns=["visitorid", "bought"]).corr()
corr.style.background_gradient(cmap="coolwarm")
```

Out[24]:

	items_view	items_addtocart	items_transaction	unique_item_views	total_sessionsduration
items_view	1.000000	0.757616	0.782903	0.978888	
items_addtocart	0.757616	1.000000	0.903854	0.735658	
items_transaction	0.782903	0.903854	1.000000	0.769944	
unique_item_views	0.978888	0.735658	0.769944	1.000000	
total_sessionsduration	0.952153	0.816382	0.864906	0.936932	
num_sessions	0.619803	0.315340	0.284208	0.572383	
mean_sessionsduration	0.222622	0.169587	0.140702	0.217528	

Nearly all of the features have high correlations (>0.7) with each other. This makes the logistic regression not usable. Another classifier that will not be used is the support vector classifier. The reasoning behind this is the sheer amount of data. The dataframe contains well over 1 million samples. A support classifier is impractical for that number of samples.[6]

In [25]:

```
nrows = procEventsDf.shape[0]
print(f"Number of Rows: {nrows}")
```

Number of Rows: 1407580

This boils it down to using the remaining four classifiers. They will be tested against each other and the one with the highest accuracy will be picked. But before doing so there are some final preparations that needs to be done.

Model Preparations

The final preparations are the following: the data will be saved in a new variable so if a mistake happens or if the model needs to be tuned a bit more it is not needed to rerun the whole notebook again. The dataset will also be limited to the range of total item views mentioned a few times above. The last step before starting to fit the models is to drop the columns `visitorid`, `items_transaction` and `items_addtocart`. The reason to drop these columns is fairly simple. `visitorid` is just a unique column in the dataset. Since it is unique for every user it doesn't generate much insights. `items_transaction` and `items_addtocart` however will make the model susceptible to overfitting. Since adding an item to a cart often leads to buying said item and if a user buys or not is the whole idea behind this modeling.

In [26]:



```

from IPython.display import Image
import pydotplus
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.externals.six import StringIO
from sklearn.tree import export_graphviz

modelDf = procEventsDf[
    (procEventsDf["items_view"] >= 3) & (procEventsDf["items_view"] <= 100)
].drop(columns=["visitorid", "items_transaction", "items_addtocart"])

```

Another thing needed to adress is the overrepresentation of non buying users. As already found out only 0.83% of every user bought something. The dataframe is already reduced. So it only holds a subset. Now lets take a look how the numbers compare to each other.

In [27]:



```

modelBought = modelDf[modelDf["bought"] == 1].shape[0]
modelBoughtPerc = modelBought / modelDf.shape[0]
modelNotBought = modelDf[modelDf["bought"] == 0].shape[0]
modelNotBoughtPerc = modelNotBought / modelDf.shape[0]

print(f"Buying Users: {modelBought} / {round(modelBoughtPerc * 100, 2)}%")
print(f"Non Buying Users: {modelNotBought} / {round(modelNotBoughtPerc * 100, 2)}%")

```

```

Buying Users: 7190 / 3.77%
Non Buying Users: 183434 / 96.23%

```

Since some are already excluded it seems the buying user have gained a portion. They now make up 3.77% of the remaining dataset. This is still dangerous to that extent as that this will definitely overfit the model if continued with the data like that. To approach this particual problem a lot of the non buying users will be excluded. This will be done by dropping some of them and run the calculation from above again.

In [28]:



```

dropRows = modelDf[modelDf["bought"] == 0].sample(frac=0.85).index.to_list()
modelDf.drop(dropRows, inplace=True)

modelBought = modelDf[modelDf["bought"] == 1].shape[0]
modelBoughtPerc = modelBought / modelDf.shape[0]
modelNotBought = modelDf[modelDf["bought"] == 0].shape[0]
modelNotBoughtPerc = modelNotBought / modelDf.shape[0]

print(f"Buying Users: {modelBought} / {round(modelBoughtPerc * 100, 2)}%")
print(f"Non Buying Users: {modelNotBought} / {round(modelNotBoughtPerc * 100, 2)}%")

```

```

Buying Users: 7190 / 20.72%
Non Buying Users: 27515 / 79.28%

```

By removing 85% of the non buying users the dataset is left with a number of 27515 users wo didn't purchase

anything. It also brought up the percentage of buying users to 20.72% of the dataset.

Fitting

After the preparations are now done the data will be split into training and test data. The test data size will be 33% of the original dataset. The following step is then fitting the different models with the training data.

In [29]:

```
X = modelDf.drop(columns=["bought"])
y = modelDf["bought"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.33, random_state=28
)

models = [
    ("Naive Bayes", GaussianNB()),
    ("Random Forest", RandomForestClassifier(max_depth=4, random_state=28)),
    ("Decision Tree", DecisionTreeClassifier(max_depth=4, random_state=28)),
    ("Nearest Neighbors", KNeighborsClassifier()),
]

for name, model in models:
    clf = model.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    accuracy = metrics.accuracy_score(y_test, y_pred)
    print(f"Accuracy for {name}: {accuracy}")
```

```
Accuracy for Naive Bayes: 0.8084344713175587
Accuracy for Random Forest: 0.8231904304549027
Accuracy for Decision Tree: 0.8245001309700515
Accuracy for Nearest Neighbors: 0.7956867196367764
```

The regular Decision Tree beat the random forest by a margin. The decision tree will therefore be evaluated in the final part.

Evaluation

To start the evaluation the confusion matrix will be utilized. This will help out to find whether the classifier just predicted the major class or had some different errors.

In [30]:



```
clf = DecisionTreeClassifier(max_depth=4, random_state=28).fit(X_train, y_train)
y_pred = clf.predict(X_test)

print(metrics.confusion_matrix(y_test, y_pred))
metrics.plot_confusion_matrix(clf, X_test, y_test)
plt.show()
```

```
[[8657  376]
 [1634  786]]
```




Figure 5: Confusion Matrix for the classifier.

The classifier tends to predict more non buying users. This can stem from a two conditions.

1. The data still has an overrepresentation of non buying users.
2. The groups don't have the clearest distinctions.

The first point should be clear. Even after breaking it down there is a ratio from 80% non buying users to 20 % buying. A bit of tweaking to ease out the overrepresentation could do the trick so that the classifier might be more able to predict better.

The second one is a bit more difficult to address. Figure 4 clearly show a lot of overlapping in between the groups. This is definitely made clear when looking at the diagonal graphs. The greatest distinction can be made in the combination of `total_sessionsduration` and `num_sessions` . But the confusion matrix and the accuracy score alone don't tell the full story. To evaluate further the precision and recall will be used. These two metrics use the confusion matrix to evaluate the performance of a model. They are calculated by the following equations:

$$Precision = \frac{TruePositive}{PredictedPositives}$$

$$Recall = \frac{TruePositives}{ActualPositives}$$

[7]

In [31]:



```
recall = metrics.recall_score(y_test, y_pred)
precision = metrics.precision_score(y_test, y_pred)

print(f"Precision Score: {precision}")
print(f"Recall Score: {recall}")
```

Precision Score: 0.6764199655765921

Recall Score: 0.3247933884297521

While the model has a good accuracy these two values tend to be fairly low. This isn't ideal since this means the predictions tend to prefer the negative group. Before diving deeper into the evaluation it will now be tested if the models performance can be boosted by reducing the overrepresentation of non buying users. A new dataframe will be used so it doesn't interfere with the old work. Another assumption made for the ongoing evaluation is whether or not the manual filtering of actions was really helpful for the model. Therefore no filtering will take place in the oncoming model tweaking. The classifier will be tested with different fractions of non buying users filtered out.



In [32]:

```
fractions = [0.96, 0.965, 0.97, 0.975, 0.98, 0.985, 0.99]
for f in fractions:
    modelDf_eval = procEventsDf.drop(
        columns=["visitorid", "items_transaction", "items_addtocart"]
    )

    dropRows = modelDf_eval[modelDf_eval["bought"] == 0].sample(frac=f).index.to_list()
    modelDf_eval.drop(dropRows, inplace=True)

    modelBought = modelDf_eval[modelDf_eval["bought"] == 1].shape[0]
    modelBoughtPerc = modelBought / modelDf_eval.shape[0]
    modelNotBought = modelDf_eval[modelDf_eval["bought"] == 0].shape[0]
    modelNotBoughtPerc = modelNotBought / modelDf_eval.shape[0]

    X = modelDf_eval.drop(columns=["bought"])
    y = modelDf_eval["bought"]

    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.33, random_state=28
    )

    clf_eval = DecisionTreeClassifier(max_depth=4, random_state=28).fit(
        X_train, y_train
    )
    y_pred = clf_eval.predict(X_test)

    accuracy = round(metrics.accuracy_score(y_test, y_pred), 3)
    precision = round(metrics.precision_score(y_test, y_pred), 3)
    recall = round(metrics.recall_score(y_test, y_pred), 3)
    print(f"Fraction: {f}")
    print(f"Buying Users: {modelBought} / {round(modelBoughtPerc * 100, 2)}%")
    print(f"Non Buying Users: {modelNotBought} / {round(modelNotBoughtPerc * 100, 2)}%")
    print(f"Accuracy Score: {round(accuracy * 100, 2)}%")
    print(f"Precision Score: {round(precision * 100, 2)}%")
    print(f"Recall Score: {round(recall * 100, 2)}%\n")
```

Fraction: 0.96
Buying Users: 11719 / 17.35%
Non Buying Users: 55834 / 82.65%
Accuracy Score: 92.1%
Precision Score: 78.4%
Recall Score: 74.9%

Fraction: 0.965
Buying Users: 11719 / 19.35%
Non Buying Users: 48855 / 80.65%
Accuracy Score: 92.0%
Precision Score: 76.8%
Recall Score: 84.1%

Fraction: 0.97
Buying Users: 11719 / 21.87%
Non Buying Users: 41876 / 78.13%
Accuracy Score: 91.3%
Precision Score: 80.7%
Recall Score: 77.5%

Fraction: 0.975

Buying Users: 11719 / 25.14%
Non Buying Users: 34897 / 74.86%
Accuracy Score: 91.5%
Precision Score: 78.7%
Recall Score: 90.6%

Fraction: 0.98
Buying Users: 11719 / 29.57%
Non Buying Users: 27917 / 70.43%
Accuracy Score: 91.9%
Precision Score: 80.3%
Recall Score: 95.6%

Fraction: 0.985
Buying Users: 11719 / 35.89%
Non Buying Users: 20938 / 64.11%
Accuracy Score: 91.5%
Precision Score: 84.9%
Recall Score: 92.8%

Fraction: 0.99
Buying Users: 11719 / 45.64%
Non Buying Users: 13959 / 54.36%
Accuracy Score: 92.7%
Precision Score: 88.4%
Recall Score: 96.6%

By not filtering out the behaviour branded as unusual the model got a direct increase in all of the values. Even if the percentages are kept at an unequal level the model gets an accuracy boost from about +10%, precision also goes up about +10% and the recall gained a plus from over 40%. A huge increase. But the best performance for all three measures is for the last fraction. This will now be used as the base model for further evaluation. The first step is now to look again at the confusion matrix.

In [33]:



```
print(metrics.confusion_matrix(y_test, y_pred))
metrics.plot_confusion_matrix(clf_eval, X_test, y_test)
plt.show()
```

```
[[4137  489]
 [ 129 3719]]
```




Figure 6: Confusion Matrix for the adjusted classifier.

As it can be seen now the reduction of the overrepresentation and the elimination of manual filtering led to better prediction results. The false predictions went down by a lot. This is a good sign that the adjustments made were indeed correct. The main part of the next step will be the feature importance.

In [34]:



```
for name, importance in zip(X.columns, clf_eval.feature_importances_):
    print(f"{name}: {round(importance, 4)}")
```

```
items_view: 0.0784
unique_item_views: 0.0
total_sessionsduration: 0.902
num_sessions: 0.0194
mean_sessionsduration: 0.0002
```

Seems like the columns `total_sessionsduration` has by far the most impact for the decisions of the tree. While the others don't have even a fraction of the importance. When compared with Figure 4 there is a slight shift to the right visible for `total_sessionsduration` in the group of buying users. If compared to the other features it's the one which seems to have the least overlap between the groups of buying users and those who don't. This was already mentioned a few times. It should therefore be no surprise that the most importance is put on this particular feature. As for now the decision process of the tree will be visualized.

In [35]:

```

dot_data = StringIO()
export_graphviz(clf_eval,
                out_file=dot_data,
                filled=True,
                rounded=True,
                special_characters=True,
                feature_names=X.columns,
                class_names=["No", "Yes"])
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())

```

Out[35]:

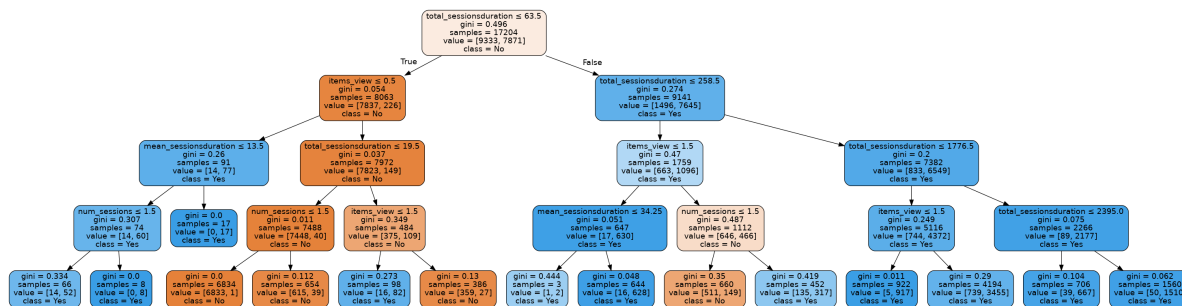


Figure 7: Decision Process of the adjusted classifier.

The tree visualization reinforces the assumptions made above. The decision trees first node tests for `total_sessionsduration` where the most positive classifications can be found on the right side, the one with the higher duration, of the tree. Overall is the testing on this feature done several times in the whole decision process.

Conclusion

The summary of the analysis will be by the lines of the tasks defined in business understanding.

1. *Explorative analysis of the data in data understanding*
2. *Cluster unusual user behaviour and filter them out*
3. *How can user behaviour be predicted?*

1. Explorative analysis of the data in data understanding

This was a feat which generated a lot of insights. While some insights can be generated from clickstream data by just looking at the raw data the most beneficial informations were generated after feature engineering and looking at the aggregated data. This was by far the most extensive part which consumed most of the work. The main key points of this task are:

- Only a small fraction of users actually buys an item.
- The greatest distinction between the groups is in the session time.

- Users who buy an item tend to have seen more items.
- Some users don't have any views on items and directly buy them.
- Non Buying users tend to have more unique sessions

2. Cluster unusual user behaviour and filter them out

This was done manually in [plotting the data by groups](#). While it helped to filter out some unusual behaviour and concentrate more on focus groups the accuracy of prediction took a big loss by it. So maybe for further analysis this shouldn't be done manually. This can be concluded by the following: it may be helpful to filter out unusual behaviour but it always comes with a loss.

3. How can user behaviour be predicted?

The model that was used was a simple classifying algorithm. Just with basic models it is possible to achieve a lot when good use of feature engineering is made. There are way more sophisticated models that could be used to predict the behaviour of a user. But just by transforming the raw data a lot of insights were generated that can help when working with raw user behaviour data like this.

Footnotes

- [1] Retailrocket (2017) Retailrocket recommender system dataset, Version 4. Retrieved 2020-04-19 from <https://www.kaggle.com/retailrocket/ecommerce-dataset> (<https://www.kaggle.com/retailrocket/ecommerce-dataset>)
- [2] Wirth, R., & Hipp, J. (2000). CRISP-DM: Towards a standard process model for data mining. In Proceedings of the 4th international conference on the practical applications of knowledge discovery and data mining (pp. 29-39). London, UK: Springer-Verlag.
- [3] Wolf Riepel (2012). CRISP-DM: Ein Standard-Prozess-Modell für Data Mining. Retrieved 2020-05-10 from <https://statistik-dresden.de/archives/1128> (<https://statistik-dresden.de/archives/1128>)
- [4] scikit-learn developers (2019). Classifier comparison. Retrieved 2020-05-30 from https://scikit-learn.org/stable/auto_examples/classification/plot_classifier_comparison.html#sphx-glr-auto-examples-classification-plot-classifier-comparison-py (https://scikit-learn.org/stable/auto_examples/classification/plot_classifier_comparison.html#sphx-glr-auto-examples-classification-plot-classifier-comparison-py)
- [5] Cohen, J., Cohen, P., West, S. G., & Aiken, L. S. (2013). Applied multiple regression/correlation analysis for the behavioral sciences. Routledge.
- [6] scikit-learn developers (2019). sklearn.svm.SVC. Retrieved 2020-05-23 from <https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html#sklearn.svm.SVC> (<https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html#sklearn.svm.SVC>)
- [7] Yousef, M., & Allmer, J. (Eds.). (2014). miRNomics: microRNA biology and computational analysis. New York, NY: Humana Press.